

Chapter 6 HW

James J. Herlan

```
grocery <- read.csv('grocery.csv')
```

```
grocery
```

	Y	X1	X2	X3
1	4264	305657	7.17	0
2	4496	328476	6.20	0
3	4317	317164	4.61	0
4	4292	366745	7.02	0
5	4945	265518	8.61	1
6	4325	301995	6.88	0
7	4110	269334	7.23	0
8	4111	267631	6.27	0
9	4161	296350	6.49	0
10	4560	277223	6.37	0
11	4401	269189	7.05	0
12	4251	277133	6.34	0
13	4222	282892	6.94	0
14	4063	306639	8.56	0
15	4343	328405	6.71	0
16	4833	321773	5.82	1
17	4453	272319	6.82	0
18	4195	293880	8.38	0
19	4394	300867	7.72	0
20	4099	296872	7.67	0
21	4816	245674	7.72	1
22	4867	211944	6.45	1
23	4114	227996	7.22	0
24	4314	248328	8.50	0
25	4289	249894	8.08	0

26	4269	302660	7.26	0
27	4347	273848	7.39	0
28	4178	245743	8.12	0
29	4333	267673	6.75	0
30	4226	256506	7.79	0
31	4121	271854	7.89	0
32	3998	293225	9.01	0
33	4475	269121	8.01	0
34	4545	322812	7.21	0
35	4016	252225	7.85	0
36	4207	261365	6.14	0
37	4148	287645	6.76	0
38	4562	289666	7.92	0
39	4146	270051	8.19	0
40	4555	265239	7.55	0
41	4365	352466	6.94	0
42	4471	426908	7.25	0
43	5045	369989	9.65	1
44	4469	472476	8.20	0
45	4408	414102	8.02	0
46	4219	302507	6.72	0
47	4211	382686	7.23	0
48	4993	442782	7.61	1
49	4309	322303	7.39	0
50	4499	290455	7.99	0
51	4186	411750	7.83	0
52	4342	292087	7.77	0

6.10

a Fit regression model (6.5) to the data for three predictor variables. State the estimated regression function. How are b_1 , b_1 , b_1 here?

```
grocery.lm <- lm(Y ~ X1 + X2 + X3, data = grocery)
```

```
grocery.lm
```

Call:

```
lm(formula = Y ~ X1 + X2 + X3, data = grocery)
```

Coefficients:

(Intercept)	X1	X2	X3
4.150e+03	7.871e-04	-1.317e+01	6.236e+02

```
summary(grocery.lm)
```

Call:

```
lm(formula = Y ~ X1 + X2 + X3, data = grocery)
```

Residuals:

Min	1Q	Median	3Q	Max
-264.05	-110.73	-22.52	79.29	295.75

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.150e+03	1.956e+02	21.220	< 2e-16 ***
X1	7.871e-04	3.646e-04	2.159	0.0359 *
X2	-1.317e+01	2.309e+01	-0.570	0.5712
X3	6.236e+02	6.264e+01	9.954	2.94e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 143.3 on 48 degrees of freedom

Multiple R-squared: 0.6883, Adjusted R-squared: 0.6689

F-statistic: 35.34 on 3 and 48 DF, p-value: 3.316e-12

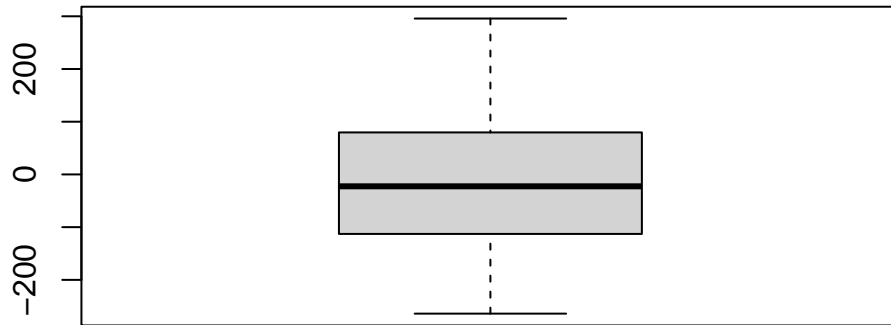
b. Obtain the residuals and prepare a box plot of the residuals. What information does this plot provide?

```
grocery.lm$residuals
```

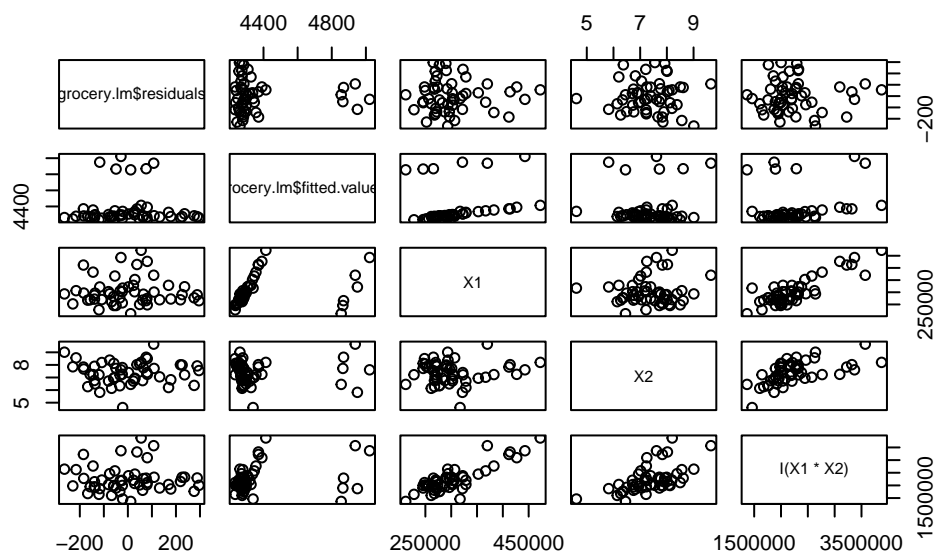
1	2	3	4	5	6
-32.063483	169.205091	-21.825426	-54.119552	75.933724	28.000660
7	8	9	10	11	12
-156.684401	-166.983381	-136.691019	275.783546	132.059842	-33.540598
13	14	15	16	17	18
-59.173782	-215.535629	22.975644	-117.076677	178.568096	-75.863154
19	20	21	22	23	24
108.947943	-183.565972	-49.165210	11.662167	-120.279732	80.569854
25	26	27	28	29	30

48.807558	-23.519661	78.869281	-58.398630	61.303251	-23.214763
31	32	33	34	35	36
-138.978271	-264.053024	218.752742	235.960794	-229.055311	-67.763118
37	38	39	40	41	42
-139.284659	288.397234	-108.609359	295.751820	29.065887	80.555515
43	44	45	46	47	48
107.399309	55.197555	37.772702	-80.508888	-144.901536	-28.753312
49	50	51	52		
2.731302	225.697849	-184.877629	64.516810		

```
boxplot(grocery.lm$residuals)
```



```
pairs(~grocery.lm$residuals+grocery.lm$fitted.values+X1+X2+I(X1*X2), data = grocery)
```

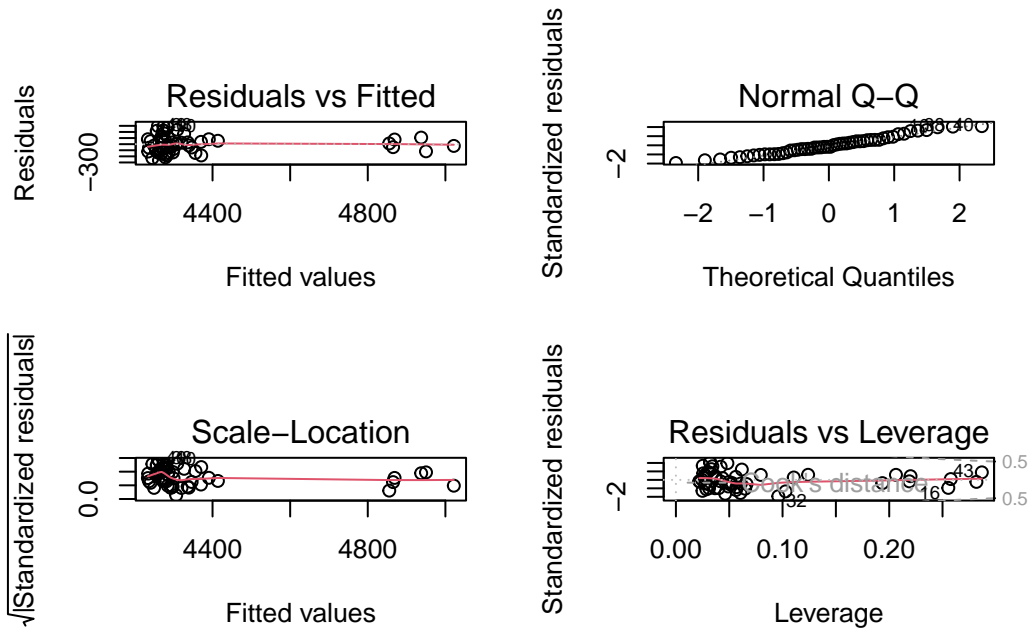


```
plot(abs(grocery.lm$residuals)~grocery.lm$fitted.values, data = grocery)
```




c. Plot the residuals against \hat{Y} , X_1 , X_2 , X_3 , and X_1X_2 on separate graphs. Also prepare a normal probability plot. Interpret the plots and summarize your findings.

```
par(mfrow = c(2, 2))
plot(grocery.lm)
```



```
qqnorm(grocery.lm$residuals)
qqline(grocery.lm$residuals)
```

Normal Q-Q Plot



```
shapiro.test(grocery.lm$residuals)
```

Shapiro-Wilk normality test

```
data:  grocery.lm$residuals  
W = 0.97575, p-value = 0.3644
```

```
lillie.test(grocery.lm$residuals)
```

Lilliefors (Kolmogorov-Smirnov) normality test

```
data:  grocery.lm$residuals  
D = 0.08161, p-value = 0.5245
```

d. Prepare a time plot of the residuals. Is there any indication that the error term are correlated?

```
grocery.resid <- resid(grocery.lm)
```

```
grocery.resid
```

1	2	3	4	5	6
-32.063483	169.205091	-21.825426	-54.119552	75.933724	28.000660
7	8	9	10	11	12
-156.684401	-166.983381	-136.691019	275.783546	132.059842	-33.540598
13	14	15	16	17	18
-59.173782	-215.535629	22.975644	-117.076677	178.568096	-75.863154
19	20	21	22	23	24
108.947943	-183.565972	-49.165210	11.662167	-120.279732	80.569854
25	26	27	28	29	30
48.807558	-23.519661	78.869281	-58.398630	61.303251	-23.214763
31	32	33	34	35	36
-138.978271	-264.053024	218.752742	235.960794	-229.055311	-67.763118
37	38	39	40	41	42
-139.284659	288.397234	-108.609359	295.751820	29.065887	80.555515
43	44	45	46	47	48
107.399309	55.197555	37.772702	-80.508888	-144.901536	-28.753312
49	50	51	52		
2.731302	225.697849	-184.877629	64.516810		

```
plot(grocery$X1, grocery.resid,
     ylab = "Residuals", xlab = "X1",
     main = "Grocery")
abline(0, 0) # the horizon
```




```
plot(grocery$X2, grocery.resid,  
     ylab = "Residuals", xlab = "X2",  
     main = "Grocery")  
abline(0, 0)           # the horizon
```



```
plot(grocery$X2, grocery.resid,  
     ylab = "Residuals", xlab = "X3",  
     main = "Grocery")  
abline(0, 0)           # the horizon
```



6.11

a. Test whether there is a regression relation, using level of significance 0.05. State the alternatives, decision rule, and conclusion. What does your result imply B_1 , B_2 , and B_3 ? What is the P -value of the test?

```
summary(grocery.lm)
```

Call:

```
lm(formula = Y ~ X1 + X2 + X3, data = grocery)
```

Residuals:

Min	1Q	Median	3Q	Max
-264.05	-110.73	-22.52	79.29	295.75

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.150e+03	1.956e+02	21.220	< 2e-16 ***
X1	7.871e-04	3.646e-04	2.159	0.0359 *
X2	-1.317e+01	2.309e+01	-0.570	0.5712
X3	6.236e+02	6.264e+01	9.954	2.94e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 143.3 on 48 degrees of freedom

Multiple R-squared: 0.6883, Adjusted R-squared: 0.6689

F-statistic: 35.34 on 3 and 48 DF, p-value: 3.316e-12

```
Anova(grocery.lm, type = 'III')
```

Anova Table (Type III tests)

Response: Y

	Sum Sq	Df	F value	Pr(>F)
(Intercept)	9245215	1	450.2861	< 2.2e-16 ***
X1	95707	1	4.6614	0.03588 *
X2	6675	1	0.3251	0.57123
X3	2034514	1	99.0905	2.941e-13 ***
Residuals	985530	48		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

$H_0: B_1 = B_2 = B_3$

H_a : not all $B_k = 0$ ($k = 1, 2, 3$).

$MSR = 725,535$

$MSE = 20,531$

$F^* = 725,535/531.9 = 35.337$

$F(0.95; 3, 48) = 2.79806$.

If $F^* < 2.79806$ conclude H_0 , otherwise H_a .

Conclude H_a . P -value = < 0.05 .

b. Estimate B_1 and B_3 jointly by the Bonferroni procedure, using a 95 percent family confidence coefficient. Interpret your result.

```
confint.lm(grocery.lm)
```

	2.5 %	97.5 %
(Intercept)	3.756677e+03	4.543098e+03
X1	5.409544e-05	1.520065e-03
X2	-5.959506e+01	3.326302e+01
X3	4.976064e+02	7.495025e+02

c. Calculate the coefficient of multiple determination R^2 . How is this measure interpreted here?

```
summary(grocery.lm)
```

Call:

```
lm(formula = Y ~ X1 + X2 + X3, data = grocery)
```

Residuals:

Min	1Q	Median	3Q	Max
-264.05	-110.73	-22.52	79.29	295.75

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.150e+03	1.956e+02	21.220	< 2e-16 ***
X1	7.871e-04	3.646e-04	2.159	0.0359 *
X2	-1.317e+01	2.309e+01	-0.570	0.5712
X3	6.236e+02	6.264e+01	9.954	2.94e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 143.3 on 48 degrees of freedom

Multiple R-squared: 0.6883, Adjusted R-squared: 0.6689

F-statistic: 35.34 on 3 and 48 DF, p-value: 3.316e-12